

Loom Video Script — Juan Jaramillo | AI Engineer

Duration: ~2:30–3:00 min | **Language:** English | **Tone:** Confident, clear, professional

[INTRO — 0:00–0:25]

[Look directly at camera, calm and confident]

“Hi — I’m Juan Jaramillo, an AI Development Lead and full-stack engineer based in Bogotá, Colombia. I have over 17 years of experience building digital products, and for the past several years I’ve been fully focused on production-grade AI systems — things that actually ship, scale, and move business metrics. I work at the intersection of generative AI, agentic workflows, and cloud-native software engineering.”

[EXPERIENCE OVERVIEW — 0:25–1:00]

“My core stack includes TypeScript, Next.js, React, Node.js, and cloud infrastructure on GCP and AWS. On the AI side, I specialize in LLMs, prompt engineering, RAG pipelines, and multi-agent orchestration using frameworks like LangChain, LangGraph, CrewAI, and Vertex AI Gemini. I’ve built production systems that serve real users — not demos, not prototypes — actual revenue-generating products deployed on Cloud Run with PostgreSQL and BigQuery backends.”

“I’ve also led AI consulting engagements for mid-market companies across the U.S. and Latin America — translating messy business problems into precise, automated systems. I don’t just design roadmaps. I own the build.”

[COMPLEX AI PROJECT — RouteGenius — 1:00–2:30]

“Now, let me walk you through one of the most technically complex AI projects I’ve built: **RouteGenius**.”

- “The problem was this: a digital media company needed to understand which traffic sources, landing pages, and monetization strategies were actually generating revenue — but their current setup had no way to test multiple variants simultaneously or attribute outcomes intelligently. They were flying blind.”
- “My approach was to build a **web traffic distribution engine** — a system that receives incoming traffic and routes it across multiple destination URLs, each assigned a specific percentage weight. Think of it as intelligent A/B routing at the infrastructure level.”
- “The real complexity came from what happened *after* routing. Each traffic segment fed into an **AI agent workflow built in n8n**, triggered via webhooks. Those agents were responsible for analyzing monetization signals in real time — evaluating click behavior, session depth, and revenue events across each route — and feeding that data back into a decision loop that could adjust traffic weights dynamically.”
- “The challenges were significant. First, webhook reliability at scale — I had to design the pipeline to handle concurrency without dropping events. Second, making the agent logic deterministic enough to produce actionable analytics without hallucinating correlations. Third, building a clean UI that let non-technical stakeholders manage route percentages and read performance data without needing to touch the underlying system.”
- “The outcome: the company gained full visibility into which traffic routes and monetization patterns were actually performing. They reduced wasted ad spend, identified their highest-converting paths, and had a repeatable framework for testing new monetization strategies — all driven autonomously by the agent layer.”

[CLOSE — 2:30–3:00]

- “That project represents how I think about AI engineering: not as a feature, but as a system. Every component — the routing logic, the agents, the data pipeline — had to work together reliably in production. That’s the standard I hold myself to.”
- “I’m excited about the opportunity at Cadre AI because you’re building exactly this kind of thing — complex, purposeful AI systems for real business impact. I’d love to bring that same mindset to your team.”
- “Thanks for watching — I’m looking forward to the conversation.”

Production Notes

Element	Recommendation
Background	Clean, minimal — desk or neutral wall
Camera	Eye level, good lighting (ring light or window light)
Pace	Slow down on technical terms — let them land

Element	Recommendation
Screen share	Optional: briefly show the RouteGenius UI at https://route.topnetworks.co/login during that section
Loom setting	Record face + screen, keep face bubble visible throughout